


Clustering time series


Jaehyoung Hong

Time-series clustering becomes important recently




Clustering : Similarity & Difference

Memory



Processor



Facebook (FB:NASDAQ)

USD

Extended Hours

Last | 4:09:22 PM EDT

160.69 +0.13 (0.081%)

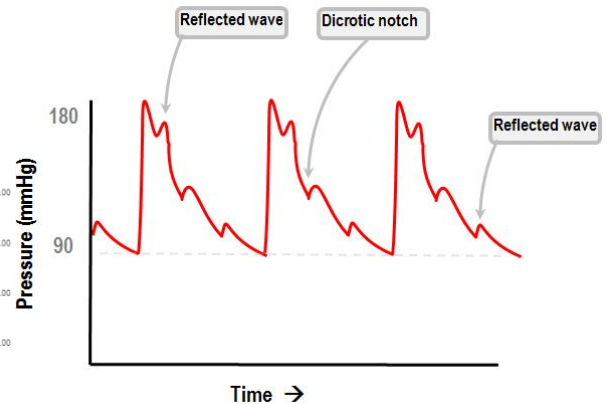
Close | 4:00:00 PM EDT

160.47 -5.51 (-3.3197%)

1 Year




CNBC




Time-series clustering needs some important characteristics

Very large data

- ECG (heart rate) : 1 (GB/hour)
- Typical weblog : 5 (GB/week)
- Space shuttle database : 200 (GB/day)





Raw data comparison is inefficient



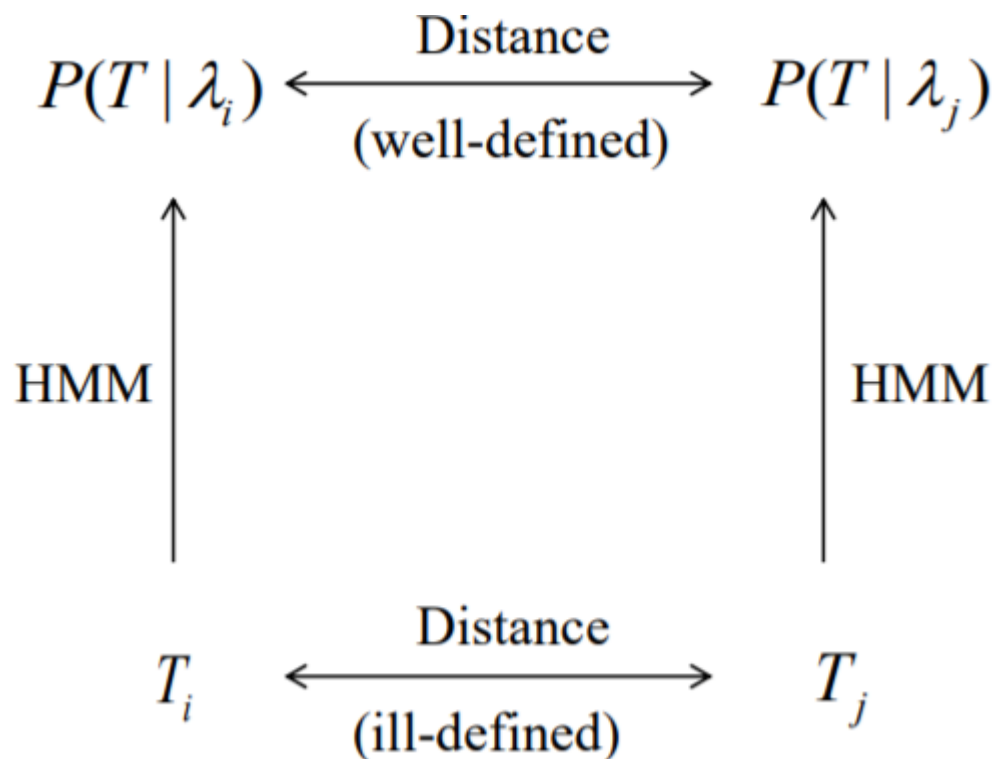
Make model / Transform to statistic

Anti-phase



Carefully chosen distance is needed

Outline of clustering with hidden markov model (HMM)




- $P(T|\lambda_i)$: Probability of observe time-series T when hidden markov model has parameter $\lambda_i = (A_i, B_i)$
- A_i : Matrix of transition probability
- B_i : Matrix of emission probability



Popular clustering method
(Partition around medoids)

Hidden markov model allows us to talk about *both observed & hidden* events



Number of ice cream (1~3; Observed) \rightarrow Weather (H or C; Hidden)

HMM based on the Markov chain

$Q = q_1 q_2 \dots q_N$	a set of N states
$A = a_{11} \dots a_{ij} \dots a_{NN}$	a transition probability matrix A , each a_{ij} representing the probability of moving from state i to state j , s.t. $\sum_{j=1}^N a_{ij} = 1 \quad \forall i$
$O = o_1 o_2 \dots o_T$	a sequence of T observations , each one drawn from a vocabulary $V = v_1, v_2, \dots, v_V$
$B = b_i(o_t)$	a sequence of observation likelihoods , also called emission probabilities , each expressing the probability of an observation o_t being generated from a state i
$\pi = \pi_1, \pi_2, \dots, \pi_N$	an initial probability distribution over states. π_i is the probability that the Markov chain will start in state i . Some states j may have $\pi_j = 0$, meaning that they cannot be initial states. Also, $\sum_{i=1}^n \pi_i = 1$

Markov Assumption: $P(q_i = a | q_1 \dots q_{i-1}) = P(q_i = a | q_{i-1})$

Output Independence: $P(o_i | q_1 \dots q_i, \dots, q_T, o_1, \dots, o_i, \dots, o_T) = P(o_i | q_i)$

Three fundamental problems of HMM is key point of clustering idea

Problem 1 (Likelihood):

Given an HMM $\lambda = (A, B)$ and an observation sequence O , determine the likelihood $P(O|\lambda)$.

<Forward Algorithm>

Problem 2 (Decoding):

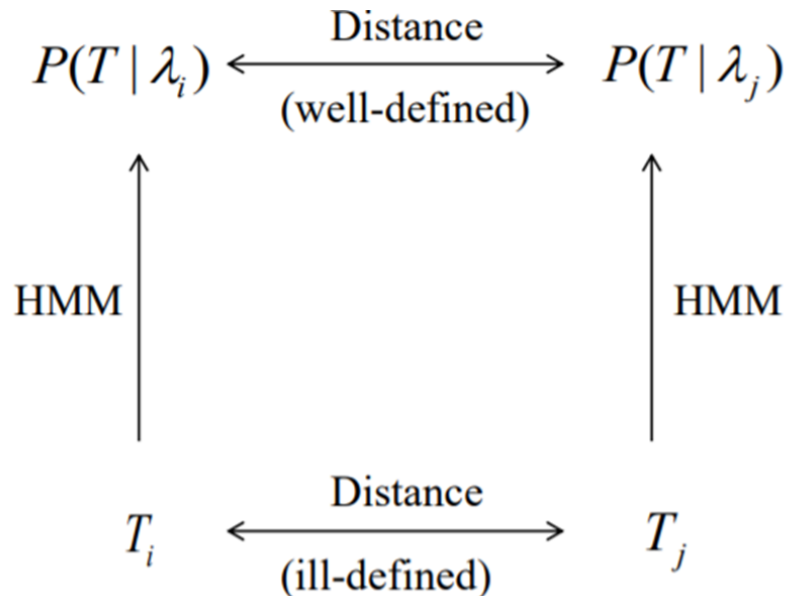
Given an observation sequence O and an HMM $\lambda = (A, B)$, discover the best hidden state sequence Q .

<Viterbi Algorithm>

Problem 3 (Learning):

Given an observation sequence O and the set of states in the HMM, learn the HMM parameters A and B .

<Forward-Backward Algorithm>



We can compute the likelihood of a particular observation by using the forward algorithm

Computing Likelihood: Given an HMM $\lambda = (A, B)$ and an observation sequence O , determine the likelihood $P(O|\lambda)$.

$$P(O|Q) = \prod_{i=1}^T P(o_i|q_i) \quad \Rightarrow \quad P(3 \ 1 \ 3 | \text{hot hot cold}) = P(3|\text{hot}) \times P(1|\text{hot}) \times P(3|\text{cold})$$


(Fully determined by emission probability)

$$P(O, Q) = P(O|Q) \times P(Q) = \prod_{i=1}^T P(o_i|q_i) \times \prod_{i=1}^T P(q_i|q_{i-1})$$
$$\Rightarrow P(3 \ 1 \ 3, \text{hot hot cold}) = P(\text{hot}|\text{start}) \times P(\text{hot}|\text{hot}) \times P(\text{cold}|\text{hot}) \times P(3|\text{hot}) \times P(1|\text{hot}) \times P(3|\text{cold})$$

(Determined by transition and emission probability)

We can compute the likelihood of a particular observation by using the forward algorithm

$$P(O) = \sum_Q P(O, Q) = \sum_Q P(O|Q)P(Q)$$

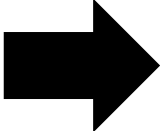


$$P(3 \ 1 \ 3) = P(3 \ 1 \ 3, \text{cold cold cold}) + P(3 \ 1 \ 3, \text{cold cold hot}) + P(3 \ 1 \ 3, \text{hot hot cold}) + \dots$$

- Problem : N^T possible hidden sequence \rightarrow Too large
- Forward algorithm : $O(N^2T)$

We can compute the likelihood of a particular observation by using the forward algorithm

$\alpha_t(j) = P(o_1, o_2 \dots o_t, q_t = j | \lambda)$: probability of being in state j after seeing the first t observations


$$\alpha_t(j) = \sum_{i=1}^N \alpha_{t-1}(i) a_{ij} b_j(o_t)$$

$\alpha_{t-1}(i)$	the previous forward path probability from the previous time step
a_{ij}	the transition probability from previous state q_i to current state q_j
$b_j(o_t)$	the state observation likelihood of the observation symbol o_t given the current state j

We can compute the likelihood of a particular observation by using the forward algorithm

function FORWARD(*observations* of len T , *state-graph* of len N) **returns** *forward-prob*

create a probability matrix *forward*[N, T]

for each state s **from** 1 **to** N **do** ; initialization step

$forward[s, 1] \leftarrow \pi_s * b_s(o_1)$

for each time step t **from** 2 **to** T **do** ; recursion step

for each state s **from** 1 **to** N **do**

$forward[s, t] \leftarrow \sum_{s'=1}^N forward[s', t-1] * a_{s', s} * b_s(o_t)$

$forwardprob \leftarrow \sum_{s=1}^N forward[s, T]$; termination step

return *forwardprob*

1. Initialization:

$$\alpha_1(j) = \pi_j b_j(o_1) \quad 1 \leq j \leq N$$

2. Recursion:

$$\alpha_t(j) = \sum_{i=1}^N \alpha_{t-1}(i) a_{ij} b_j(o_t); \quad 1 \leq j \leq N, 1 < t \leq T$$

3. Termination:


$$P(O|\lambda) = \sum_{i=1}^N \alpha_T(i)$$

We can find most probable sequence of hidden states by using Viterbi algorithm

Decoding: Given as input an HMM $\lambda = (A, B)$ and a sequence of observations $O = o_1, o_2, \dots, o_T$, find the most probable sequence of states $Q = q_1 q_2 q_3 \dots q_T$.

- Naïve method : For all possible hidden sequence, compute likelihood

$$v_t(j) = \max_{q_1, \dots, q_{t-1}} P(q_1 \dots q_{t-1}, o_1, o_2 \dots o_t, q_t = j | \lambda)$$



$$v_t(j) = \max_{i=1}^N v_{t-1}(i) a_{ij} b_j(o_t)$$

We can find most probable sequence of hidden states by using Viterbi algorithm

function VITERBI(*observations* of len T , *state-graph* of len N) **returns** *best-path*, *path-prob*

create a path probability matrix *viterbi*[N , T]

for each state s **from** 1 **to** N **do** ; initialization step

$viterbi[s,1] \leftarrow \pi_s * b_s(o_1)$

$backpointer[s,1] \leftarrow 0$

for each time step t **from** 2 **to** T **do** ; recursion step

for each state s **from** 1 **to** N **do**

$viterbi[s,t] \leftarrow \max_{s'=1}^N viterbi[s',t-1] * a_{s',s} * b_s(o_t)$

$backpointer[s,t] \leftarrow \operatorname{argmax}_{s'=1}^N viterbi[s',t-1] * a_{s',s} * b_s(o_t)$

$bestpathprob \leftarrow \max_{s=1}^N viterbi[s,T]$; termination step

$bestpathpointer \leftarrow \operatorname{argmax}_{s=1}^N viterbi[s,T]$; termination step

$bestpath \leftarrow$ the path starting at state $bestpathpointer$, that follows backpointer

return $bestpath$, $bestpathprob$

Finally, we can give a formal definition of the Viterbi recursion as follows:

1. Initialization:

$$v_1(j) = \pi_j b_j(o_1) \quad 1 \leq j \leq N$$

$$bt_1(j) = 0 \quad 1 \leq j \leq N$$

2. Recursion

$$v_t(j) = \max_{i=1}^N v_{t-1}(i) a_{ij} b_j(o_t); \quad 1 \leq j \leq N, 1 < t \leq T$$

$$bt_t(j) = \operatorname{argmax}_{i=1}^N v_{t-1}(i) a_{ij} b_j(o_t); \quad 1 \leq j \leq N, 1 < t \leq T$$

3. Termination:

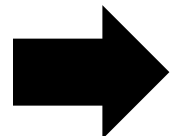
$$\text{The best score: } P^* = \max_{i=1}^N v_T(i)$$

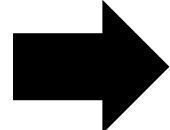
$$\text{The start of backtrace: } q_T^* = \operatorname{argmax}_{i=1}^N v_T(i)$$

We can learn the parameters of HMM by using forward-backward algorithm

Learning: Given an observation sequence O and the set of possible states in the HMM, learn the HMM parameters A and B .

3	3	2	1	1	2	1	2	3
hot	hot	cold	cold	cold	cold	cold	hot	hot


$$\pi_h = 1/3 \quad \pi_c = 2/3$$


$$\begin{array}{llll} p(\text{hot}|\text{hot}) = 2/3 & p(\text{cold}|\text{hot}) = 1/3 & P(1|\text{hot}) = 0/4 = 0 & p(1|\text{cold}) = 3/5 = .6 \\ p(\text{cold}|\text{cold}) = 2/3 & p(\text{hot}|\text{cold}) = 1/3 & P(2|\text{hot}) = 1/4 = .25 & p(2|\text{cold}) = 2/5 = .4 \\ & & P(3|\text{hot}) = 3/4 = .75 & p(3|\text{cold}) = 0 \end{array}$$

<A : transition prob>

<B : emission prob>

We can learn the parameters of HMM by using forward-backward algorithm

$\beta_t(i) = P(o_{t+1}, o_{t+2} \dots o_T | q_t = i, \lambda)$: backward probability

1. Initialization:

$$\beta_T(i) = 1, \quad 1 \leq i \leq N$$

2. Recursion

$$\beta_t(i) = \sum_{j=1}^N a_{ij} b_j(o_{t+1}) \beta_{t+1}(j), \quad 1 \leq i \leq N, 1 \leq t < T$$

3. Termination:

$$P(O|\lambda) = \sum_{j=1}^N \pi_j b_j(o_1) \beta_1(j)$$

We can learn the parameters of HMM by using forward-backward algorithm

$$\hat{a}_{ij} = \frac{\text{expected number of transitions from state } i \text{ to state } j}{\text{expected number of transitions from state } i}$$

➡ $\xi_t(i, j) = P(q_t = i, q_{t+1} = j | O, \lambda)$ (summation of it gives numerator)

➡ not-quite- $\xi_t(i, j) = P(q_t = i, q_{t+1} = j, O | \lambda)$

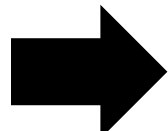
not-quite- $\xi_t(i, j) = \alpha_t(i) a_{ij} b_j(o_{t+1}) \beta_{t+1}(j)$

➡ $P(O | \lambda) = \sum_{j=1}^N \alpha_t(j) \beta_t(j)$ ➡ $\xi_t(i, j) = \frac{\alpha_t(i) a_{ij} b_j(o_{t+1}) \beta_{t+1}(j)}{\sum_{j=1}^N \alpha_t(j) \beta_t(j)}$

We can learn the parameters of HMM by using forward-backward algorithm

$$\hat{a}_{ij} = \frac{\text{expected number of transitions from state } i \text{ to state } j}{\text{expected number of transitions from state } i}$$

$$\xi_t(i, j) = \frac{\alpha_t(i) a_{ij} b_j(o_{t+1}) \beta_{t+1}(j)}{\sum_{j=1}^N \alpha_t(j) \beta_t(j)}$$



$$\hat{a}_{ij} = \frac{\sum_{t=1}^{T-1} \xi_t(i, j)}{\sum_{t=1}^{T-1} \sum_{k=1}^N \xi_t(i, k)}$$

We can learn the parameters of HMM by using forward-backward algorithm

<recomputing observation probability>

$$\hat{b}_j(v_k) = \frac{\text{expected number of times in state } j \text{ and observing symbol } v_k}{\text{expected number of times in state } j}$$

$$\gamma_t(j) = P(q_t = j | O, \lambda) \quad (\text{summation of it gives numerator})$$



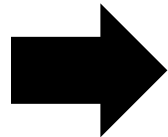
$$\gamma_t(j) = \frac{P(q_t = j, O | \lambda)}{P(O | \lambda)}$$

$$\gamma_t(j) = \frac{\alpha_t(j)\beta_t(j)}{P(O | \lambda)}$$

We can learn the parameters of HMM by using forward-backward algorithm

<recomputing observation probability>

$$\hat{b}_j(v_k) = \frac{\text{expected number of times in state } j \text{ and observing symbol } v_k}{\text{expected number of times in state } j}$$


$$\hat{b}_j(v_k) = \frac{\sum_{t=1}^T \mathbb{1}_{\{O_t=v_k\}} \gamma_t(j)}{\sum_{t=1}^T \gamma_t(j)}$$

We can learn the parameters of HMM by using forward-backward algorithm

function FORWARD-BACKWARD(*observations of len T, output vocabulary V, hidden state set Q*) **returns** HMM=(A,B)

initialize A and B

iterate until convergence

E-step

$$\gamma_t(j) = \frac{\alpha_t(j)\beta_t(j)}{\alpha_T(q_F)} \quad \forall t \text{ and } j$$

$$\xi_t(i, j) = \frac{\alpha_t(i) a_{ij} b_j(o_{t+1}) \beta_{t+1}(j)}{\alpha_T(q_F)} \quad \forall t, i, \text{ and } j$$

M-step

$$\hat{a}_{ij} = \frac{\sum_{t=1}^{T-1} \xi_t(i, j)}{\sum_{t=1}^{T-1} \sum_{k=1}^N \xi_t(i, k)}$$


$$\hat{b}_j(v_k) = \frac{\sum_{t=1 \text{ s.t. } O_t=v_k}^T \gamma_t(j)}{\sum_{t=1}^T \gamma_t(j)}$$

return A, B

Define distance for clustering using KL-divergence


- Learning : find $\lambda = (A, B)$

- Likelihood : find $P(T|\lambda)$



KL-divergence

$$D_{KL}(P(T | \lambda_i); P(T | \lambda_j)) \equiv \int dT P(T | \lambda_i) \log \frac{P(T | \lambda_i)}{P(T | \lambda_j)}$$



Monte-Carlo

$$D_{KL}(P(T | \lambda_i); P(T | \lambda_j)) \approx \frac{1}{n} \sum_{\alpha=1}^N \log \frac{P(T_{\alpha} | \lambda_i)}{P(T_{\alpha} | \lambda_j)}$$

Need large N

One point
approximation

$$D_{KL}(P(T | \lambda_i); P(T | \lambda_j)) \approx \log \frac{P(T_i | \lambda_i)}{P(T_i | \lambda_j)}$$


Too extreme

Define distance for clustering using KL-divergence

- The observed data set is sufficiently representative of the universe of possible trajectories

$$P(T | \lambda) \rightarrow \tilde{P}_\lambda \equiv \frac{1}{Z_\lambda} \{P(T_1 | \lambda), P(T_2 | \lambda), \dots, P(T_N | \lambda)\} \equiv \left\{ \tilde{P}(T_1 | \lambda), \tilde{P}(T_2 | \lambda), \dots, \tilde{P}(T_N | \lambda) \right\}$$

Where $Z_\lambda = \sum_{i=1}^N P(T_i | \lambda)$



$$D_{KL}(\lambda_i; \lambda_j) \equiv D_{KL}(\tilde{P}_{\lambda_i}; \tilde{P}_{\lambda_j}) = \sum_{i=1}^N \tilde{P}(T_i | \lambda_i) \log \frac{\tilde{P}(T_i | \lambda_i)}{\tilde{P}(T_i | \lambda_j)}$$

$$D_{ij} \equiv D(T_i; T_j) \equiv \frac{1}{2} (D_{KL}(\lambda_i; \lambda_j) + D_{KL}(\lambda_j; \lambda_i))$$

Partitioning around medoids (PAM) is clustering method similar to k-means clustering

- K-means clustering : Choose K-center → Assign data point to nearest center
→ Recompute centers as mean of data points in each cluster
- Partitioning around medoids (PAM) : Medoids must be the data points in cluster

Application to 1225 disease data shows 4 clusters

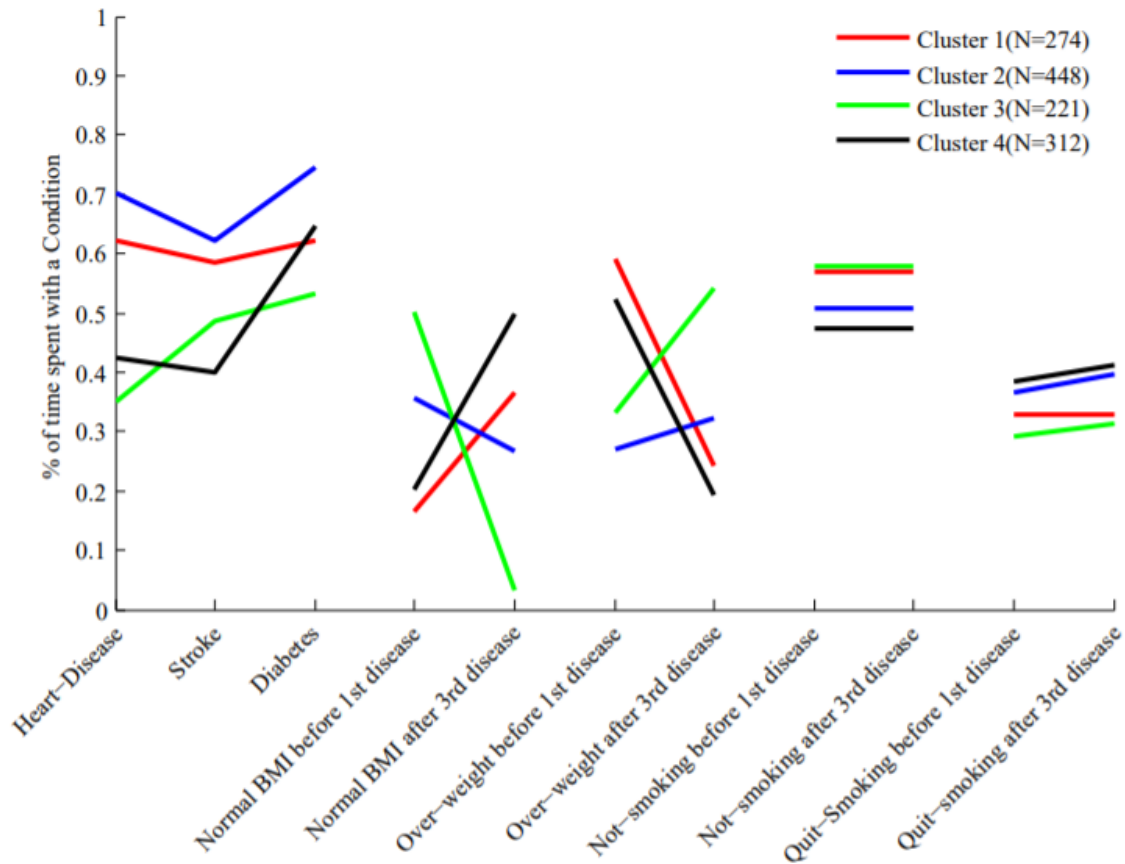
- 1225 patient = time-series / length = 18 / all patients develops three chronic condition / using obesity and smoking as covariates of hidden sequence

	Disease onset	Normal BMI	Overweight	Quit smoking
Cluster 1	Heart disease, stroke and diabetes almost simultaneously	Mostly overweight/obese before 1st disease	Some weight loss after 3rd disease	No change in smoking behaviour after 3rd disease
Cluster 2	Diabetes, heart disease and then stroke	Significantly overweight/obese before 1st disease	Some weight gain after 3rd disease	Mild increase in quitting smoking after 3rd disease
Cluster 3	Diabetes, stroke and then heart disease	Half time normal BMI before 1st disease	Significant weight gain after 3rd disease	Mild increase in quitting after 3rd disease
Cluster 4	Diabetes and then heart disease and stroke	Mostly overweight/obese before 1st disease	Significant weight loss after 3rd disease	Mild increase in quitting after 3rd disease

Number of clusters are chosen by DB / Dunn index

Application to 1225 disease data shows 4 clusters

- 1225 patient = time-series / length = 18 / all patients develops three chronic condition / using obesity and smoking as covariates of hidden sequence



Number of Hidden states	Number of clusters	Correlation
2 Hidden states	$K = 2$	0.74
3 Hidden states	$K = 4$	0.41
4 Hidden states	$K = 3$	0.47